



Rust with OpenTelemetry

Tirslen Pather


Discord: @tirslo

GitHub: Tirslo

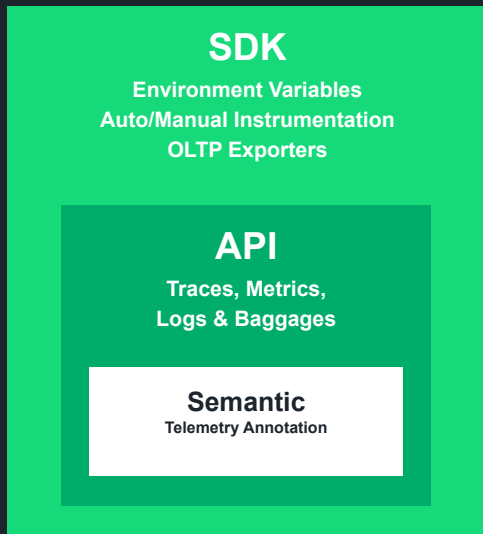


Key Facts on OpenTelemetry

- **OpenTelemetry** is a member of CNCF.
- Formed through a merger of the **OpenTracing** and **OpenCensus** projects.
- Vendor agnostic - set of APIs, libraries, integrations, and a collector service for telemetry.
- **Standardizes** how you collect telemetry data from your applications and services.
- Send it to an Observability platform of your choice.



OpenTelemetry Instrumentation



Core Concepts on Instrumentation

- **Semantic Conventions** - annotate telemetry with attributes specific to the represented operation, such as HTTP calls.
- **API** - data types for tracing, metrics, and logging data.
- **SDK** - language-specific implementation of the API.
- SDKs incorporate automatic instrumentation for common libraries and frameworks for your application.
- **OpenTelemetry Protocol (OTLP)** - used to send data to your backend Observability platform of choice.

Documentation???

Docs

What is OpenTelemetry?

▶ Getting Started

▶ Concepts

▶ Demo

▼ Instrumentation

▶ C++

▶ .NET

▶ Erlang/Elixir

▶ Go

▶ Java

▼ JavaScript

▼ Getting Started

Node.js

Browser

▶ Automatic

Manual

Libraries

Exporters

Context

Propagation

Resources

Sampling

Serverless

API

Examples

▶ PHP

▶ Python

▶ Ruby

▶ Rust

▶ Swift

Other

▶ Collector

▶ K8s Operator

▶ FaaS

▶ Migration

▶ Specs

Acknowledgements

The following OpenTelemetry libraries, exporters, and auto-instrumentation modules are available for OpenTelemetry.

Dependencies

First, install the Node SDK and autoinstrumentations package.

The Node SDK lets you initialize OpenTelemetry with several configuration defaults that are correct for the majority of use cases.

The `auto-instrumentations-node` package installs instrumentation packages that will automatically create spans corresponding to code called in libraries. In this case, it provides instrumentation for Express, letting the example app automatically create spans for each incoming request.

```
npm install @opentelemetry/sdk-node \
  @opentelemetry/api \
  @opentelemetry/auto-instrumentations-node \
  @opentelemetry/sdk-metrics
```

To find all autoinstrumentation modules, you can look at the [registry](#).

Setup

The instrumentation setup and configuration must be run *before* your application code. One tool commonly used for this task is the `require` flag.

Create a file named `instrumentation.ts` (or `instrumentation.js` if not using typescript), which will contain your instrumentation setup code.

TypeScript JavaScript

```
/**instrumentation.ts*/
import { NodeSDK } from '@opentelemetry/sdk-node';
import { ConsoleSpanExporter } from '@opentelemetry/sdk-trace-node';
import { getNodeAutoInstrumentations } from '@opentelemetry/auto-instrumentations-node';
import { PeriodicExportingMetricReader, ConsoleMetricExporter } from '@opentelemetry/sdk-metrics';

const sdk = new NodeSDK({
  traceExporter: new ConsoleSpanExporter(),
  metricReader: new PeriodicExportingMetricReader({
    exporter: new ConsoleMetricExporter()
  }),
  autoInstrumentations: [getNodeAutoInstrumentations()]
});

sdk
  .start()
```

OpenTelemetry is split into two parts: an API to instrument code with API. To start integrating OpenTelemetry into any project, the API is generated. To generate tracing telemetry in your application you will API from the [go.opentelemetry.io/otel/trace](#) package.

First, you need to install the necessary packages for the Trace API. Run your working directory.

```
go get go.opentelemetry.io/otel \
  go.opentelemetry.io/otel/trace
```

Now that the packages installed you can start updating your application the `app.go` file.

```
import (
    "context"
    "fmt"
    "io"
    "log"
    "strconv"

    "go.opentelemetry.io/otel"
    "go.opentelemetry.io/otel/attribute"
    "go.opentelemetry.io/otel/trace"
)
```

With the imports added, you can start instrumenting.

The OpenTelemetry Tracing API provides a [Tracer](#) to create traces to be associated with one instrumentation library. That way telemetry understood to come from that part of a code base. To uniquely identify a [Tracer](#) you will create a constant with the package name in `app.go`.

```
// name is the Tracer name used to identify this instrumentation
const name = "fib"
```

Using the full-qualified package name, something that should be a standard way to identify a [Tracer](#). If your example package name name you use here to match.

Docs

What is OpenTelemetry?

▶ Getting Started

▶ Concepts

▶ Demo

▼ Instrumentation

▶ C++

▶ .NET

▶ Erlang/Elixir

▶ Go

▼ Java

Getting Started

▶ Automatic

Manual

API

Examples

▶ JavaScript

▶ PHP

▶ Python

▶ Ruby

▶ Rust

▶ Swift

Other

▶ Collector

▶ K8s Operator

▶ FaaS

▶ Migration

▶ Specs

Acknowledgements

Instrumentation

Next, you'll use a [Java agent](#) to automatically instrument the application at launch time. While you can [configure the Java agent](#) in a number of ways, the steps below use environment variables.

1. Download [opentelemetry-javaagent.jar](#) from [Releases](#) of the [opentelemetry-java-instrumentation](#) repo. The JAR file contains the agent and all automatic instrumentation packages:

```
curl -L -O https://github.com/open-telemetry/opentelemetry-java-instrumentation/releases
```

Take note of the path to the JAR file.

2. Set and export variables that specify the Java agent JAR and a [console exporter](#), using a notation suitable for your shell/terminal environment — we illustrate a notation for bash-like shells:

```
$ export JAVA_TOOL_OPTIONS="-javaagent:PATH/TO/opentelemetry-javaagent.jar" \
  OTEL_TRACES_EXPORTER=logging \
  OTEL_METRICS_EXPORTER=logging \
  OTEL_LOGS_EXPORTER=logging
```

Important

Replace `PATH/TO` above, with your path to the JAR.

3. Run your **application** once again:

```
$ java -jar ./build/libs/java-simple.jar
...
```

Note the output from the `otel.javaagent`.

4. From *another* terminal, send a request using `curl`:

```
$ curl localhost:8080/rolldice
```

Docs

What is OpenTelemetry?

► Getting Started

► Concepts

► Demo

▼ Instrumentation

► C++

► .NET

► Erlang/Elixir

► Go

► Java

► JavaScript

► PHP

► Python

► Ruby

► Rust

► Swift

Other

► Collector

► K8s Operator

► FaaS

► Migration

► Specs

Acknowledgements

[Docs](#) / [Instrumentation](#) / Rust

Rust

 A language-specific implementation of OpenTelemetry in Rust.

This is the OpenTelemetry Rust documentation. OpenTelemetry is an observability framework – an API, SDK, and tools that are designed to aid in the generation and collection of application telemetry data such as metrics, logs, and traces. This documentation is designed to help you understand how to get started using OpenTelemetry Rust.

Status and Releases

The current status of the major functional components for OpenTelemetry Rust is as follows:

Traces Metrics Logs

Beta

Alpha

Not yet implemented

For releases, including the [latest release](#), see [Releases](#).


Crates

OpenTelemetry for Rust publishes the following crates:

- [opentelemetry](#)
- [opentelemetry-api](#)
- [opentelemetry-sdk](#)
- [opentelemetry-aws](#)
- [opentelemetry-contrib](#)
- [opentelemetry-datadog](#)
- [opentelemetry-dynatrace](#)
- [opentelemetry-http](#)
- [opentelemetry-jaeger](#)
- [opentelemetry-otlp](#)
- [opentelemetry-prometheus](#)
- [opentelemetry-semantic-conventions](#)
- [opentelemetry-stackdriver](#)
- [opentelemetry-zipkin](#)

Further Reading

- [Docs for Rust API & SDK](#)
- [Examples](#)
- [Ecosystem](#)



Rust App

```
use actix_web::{web, App, HttpResponse, HttpServer};
use rand::Rng;
use std::io::{self};
use std::thread;
use std::time::Duration;

async fn hello() → HttpResponse {
    HttpResponse::Ok().body("Hello World!")
}

async fn index() → HttpResponse {
    let rand_num = gen_number();
    let resp_body = format!("Random number: {}", rand_num.to_string());
    HttpResponse::Ok().body(resp_body)
}

fn gen_number() → u32 {
    let mut rng = rand::thread_rng();
    let delay = rng.gen_range(0..=5);
    let duration = Duration::from_secs(delay);
    thread::sleep(duration);

    rng.gen()
}

#[actix_rt::main]
async fn main() → io::Result<()> {
    HttpServer::new(move || {
        App::new()
            .route("/", web::get().to(hello))
            .route("/index", web::get().to(index))
    })
    .bind("127.0.0.1:8080")?
    .run()
    .await?;

    Ok(())
}
```

snappify.com

```
[package]
name = "meetup-basic"
version = "0.1.0"
edition = "2021"
```

```
[dependencies]
actix-web = "4.3.1"
actix-rt = "2.4"
rand = "0.8.5"
```

snappify.com

What's needed for Otel

```
[dependencies]
actix-web = "4.3.1"
actix-rt = "2.4"
actix-web-opentelemetry = { version = "0.13", features = ["metrics"] }
opentelemetry = { version = "0.19", features = ["metrics", "rt-tokio-current-thread", "rt-tokio"] }
opentelemetry_api = {version= "0.19.0", features = ["metrics"]}
opentelemetry_sdk = {version= "0.19.0", features = ["rt-tokio"]}
opentelemetry-otlp = { version = "0.12", features = ["http-proto", "request-client"] }
opentelemetry-semantic-conventions = "0.11.0"
request = {version = "0.11.13", features = ["blocking", "json"]}
tokio = {version = "1.0", features = ["full"]}
```

Adding Instrumentation → Tracer

```
use opentelemetry::global;
use opentelemetry::sdk::{trace, Resource};
use opentelemetry::trace::{FutureExt, TraceContextExt, Tracer};
use opentelemetry::Key;
use opentelemetry::{
    global::shutdown_tracer_provider, sdk::trace as sdktrace, trace::TraceError, KeyValue,
};
use opentelemetry_otlp::WithExportConfig;
use std::collections::HashMap;

fn init_tracer() → Result<sdktrace::Tracer, TraceError> {
    opentelemetry_otlp::new_pipeline()
        .tracing()
        .with_exporter(
            opentelemetry_otlp::new_exporter()
                .http()
                .with_endpoint("https://otlp.nr-data.net/v1/traces")
                .with_headers(HashMap::from([
                    ("api-key".to_string(), " ".to_string()),
                ]))
                .with_timeout(std::time::Duration::from_secs(2)),
        )
        .with_trace_config(
            trace::config().with_resource(Resource::new(vec![KeyValue::new(
                "service.name",
                "rust-meetup",
            )])),
        )
        .install_batch(opentelemetry::runtime::Tokio)
}
```

Initialise an OTLP Pipeline

endpoint for exporter

configuration options

makes the tracer ready
for use in the application


```

async fn index() → HttpResponse {
    let tracer = global::tracer("request");

    let rand_num = gen_number();
    let resp_body = format!("Random number: {}", rand_num.to_string());
    tracer.in_span("index", |ctx| {
        ctx.span()
            .set_attribute(Key::new("parameter").i64(rand_num.into()));
    });
    HttpResponse::Ok().body(resp_body)
}

fn gen_number() → u32 {
    let mut rng = rand::thread_rng();
    let delay = rng.gen_range(0...5);
    let duration = Duration::from_secs(delay);
    thread::sleep(duration);

    rng.gen()
}

#[actix_rt::main]
async fn main() → io::Result<()> {
    let tracer = init_tracer().expect("Failed to initialise tracer.");

    HttpServer::new(move || {
        let tracer = tracer.clone();
        App::new()
            .wrap_fn(move |req, srv| {
                tracer.in_span("middleware", move |cx| {
                    cx.span()
                        .set_attribute(Key::new("path").string(req.path().to_string()));
                    srv.call(req).with_context(cx)
                })
            })
            .route("/", web::get().to(hello))
            .route("/index", web::get().to(index))
    })
    .bind("127.0.0.1:8080")?
    .run()
    .await?;

    shutdown_tracer_provider();

    Ok(())
}

```

Create a request tracer

Create a span called index
Decorate the span with an
attribute named
parameter and value of rand_num

initialise a tracer

start each request with
a middleware span

Shutdown the provider and
export all remaining spans

What's it look like?



Tracing API

```
use actix_web::{web, App, HttpResponse, HttpServer};
use actix_web_opentelemetry::RequestTracing;
use opentelemetry::sdk::{trace, Resource};
use opentelemetry::sdk::trace as sdktrace, trace::TraceError, KeyValue;
use opentelemetry_otlp::WithExportConfig;
```

```
use rand::Rng;
use std::collections::HashMap;
use std::io;
use std::thread;
use std::time::Duration;
```

```
use tracing::{error, info, info_span, warn};
use tracing_subscriber::layer::SubscriberExt;
```

```
fn init_tracer() -> Result<sdktrace::Tracer, TraceError> {
    ....the same
}
```

```
#[tracing::instrument]
async fn hello() -> HttpResponse {
    let root_span = info_span!("Hello");
    let _entered = root_span.enter();
    info!(message = "hello_gen", "generating number");
    let rand_num = gen_number();
    index().await;
    HttpResponse::Ok().body("Hello World!")
}
```

```
.....

#[actix_rt::main]
async fn main() -> io::Result<()> {
    let tracer = init_tracer().unwrap();
    let telemetry = tracing_opentelemetry::layer().with_tracer(tracer);
    let subscriber = tracing_subscriber::Registry::default().with(telemetry);
    tracing::subscriber::set_global_default(subscriber).unwrap();
```

```
    HttpServer::new(move || {
        App::new()
            .wrap(RequestTracing::new())
            .route("/", web::get().to(hello))
            .route("/index", web::get().to(index))
    })
    .bind("127.0.0.1:8080")?
    .run()
    .await
}
```

tracing api crates

tracing macro

Configuration of span details much simpler

initialising a tracing subscriber

hello

rust-meetup

Showing the slowest spans. Use the search filter to find specific spans.

View by

Trace

5.5 s

5000 ms

4500 ms

4000 ms

3500 ms

3000 ms

2500 ms

2000 ms

1500 ms

1000 ms

500 ms

0 s

hello (rust-meetup)

Plot

Trace

hello

rust-meetup

rust-meetup

June 13 at 5:45pm

Trace ID: 1bf94558d38ff1a13e8c4bdfae978f94

Spans (5)

Trace details

Logs (0)

0 anomalous spans

0 errors

1 entity

Find spans by name or ID

Maximize

0 ms

600 ms

1200 ms

1800 ms

2401 ms

3001 ms

Expand all

Collapse all

Standard

Manual

Reset

4

hello

rust-meetup

3.00 s

Show in-process spans (4)

</> Hello

3.00 s

</> gen_number

0.03 ms

</> index

3.00 s

</> gen_number

3.00 s

Trace duration

3000.74 ms

Back-end duration

3000.74 ms

View span events

gen_number

rust-meetup

3.00 s

DURATION

Performance

Attributes

Details

Average duration

4000 ms

3000 ms

2000 ms

1000 ms

0 ms

5:30pm

5:40pm

5:50pm

Response Time

Span event

Throughput (rpm)

15

10

5

0

5:30pm

5:40pm

5:50pm

Throughput

Span event

Span duration histogram

8

6

4

2

0

0 ms

1000 ms

2000 ms

3000 ms

Duration

Documentation

- <https://tokio.rs/tokio/topics/tracing>
- <https://opentelemetry.io/docs/instrumentation/rust>
- <https://github.com/open-telemetry/opentelemetry-rust>