

Trust the Compiler

The cool things the compiler does behind your back
(and how you can make use of them)

The compiler does a LOT

Rust and LLVM each have complex optimization passes

Things like data flow and control flow analysis optimizes so many simple things

Conditional flow

Consider this simple function:

```
#[no_mangle]
pub fn run_conditionally(val: u32) {
    let reduced = val % 2;

    match reduced {
        0 => foo(),
        1 => bar(),
        _ => unreachable_fn(),
    }
}
```

Compiled LLVM IR

The unreachable code gets optimized away

```
define void @run_conditionally(i32 noundef %val) unnamed_addr #1 {  
    %0 = and i32 %val, 1  
    %trunc.not = icmp eq i32 %0, 0  
    br i1 %trunc.not, label %bb2, label %bb3  
  
bb2:                                     ; preds = %start  
    tail call void @foo()  
    br label %bb4  
  
bb3:                                     ; preds = %start  
    tail call void @bar()  
    br label %bb4  
  
bb4:                                     ; preds = %bb3, %bb2  
    ret void  
}
```

This can be taken further...

Assume the below function. What if `index` is always < 10 ?

```
pub fn index_under_10(arr: [i32; 10], index: usize) -> i32 {  
    arr[index]  
}
```

Rust still generates the check, and a panic call

But you still want to force it away

Instead of resorting to raw pointer operations, simply add an assertion

```
pub fn index_under_10_no_panic(arr: [i32; 10], index: usize) -> i32 {  
    unsafe {  
        if index >= 10 {  
            std::hint::unreachable_unchecked();  
        }  
    }  
    arr[index]  
}
```

And it's gone!

```
example::index_under_10:
    mov     rax, rsi
    cmp     rsi, 9
    ja      .LBB0_2
    mov     eax, dword ptr [rdi + 4*rax]
    ret

.LBB0_2:
    push    rax
    lea     rdx, [rip + .L_unnamed_1]
    mov     esi, 10
    mov     rdi, rax
    call    qword ptr [rip + core::panicking::panic_bounds_check@GOTPCREL]
    ud2

example::index_under_10_no_panic:
    mov     eax, dword ptr [rdi + 4*rsi]
    ret
```

Static vs Dynamic arrays

Let's say you have an array that's always only 4 elements long.

You need to perform an action on it. Say, sum:

```
#[inline(never)]  
pub fn sum_vec(arr: Vec<f32>) -> f32 {  
    arr.iter().sum()  
}
```


What actually gets generated?

Let's see the assembly for this simple function

That's a lot of instructions...

For just 4 elements

```
example::sum_vec:
    push    rax
    mov     rax, rdi
    mov     rdi, qword ptr [rdi]
    mov     rsi, qword ptr [rax + 16]
    test    rsi, rsi
    je      .LBB0_1
    mov     ecx, esi
    and     ecx, 7
    cmp     rsi, 8
    jae     .LBB0_4
    xorps   xmm0, xmm0
    xor     edx, edx
    jmp     .LBB0_6

.LBB0_1:
    xorps   xmm0, xmm0
    jmp     .LBB0_9

.LBB0_4:
    and     rsi, -8
    xorps   xmm0, xmm0
    xor     edx, edx

.LBB0_5:
    addss   xmm0, dword ptr [rdi + 4*rdx]
    addss   xmm0, dword ptr [rdi + 4*rdx + 4]
    addss   xmm0, dword ptr [rdi + 4*rdx + 8]
    addss   xmm0, dword ptr [rdi + 4*rdx + 12]
    addss   xmm0, dword ptr [rdi + 4*rdx + 16]
    addss   xmm0, dword ptr [rdi + 4*rdx + 20]
    addss   xmm0, dword ptr [rdi + 4*rdx + 24]
    addss   xmm0, dword ptr [rdi + 4*rdx + 28]
    add     rdx, 8
    cmp     rsi, rdx
    jne     .LBB0_5

.LBB0_6:
    test    rcx, rcx
    je      .LBB0_9
    lea     rdx, [rdi + 4*rdx]
    xor     esi, esi

.LBB0_8:
    addss   xmm0, dword ptr [rdx + 4*rsi]
    inc     rsi
    cmp     rcx, rsi
    jne     .LBB0_8

.LBB0_9:
    mov     rsi, qword ptr [rax + 8]
    test    rsi, rsi
    je      .LBB0_11
    shl     rsi, 2
    mov     edx, 4
    movss   dword ptr [rsp + 4], xmm0
    call    qword ptr [rip + __rust_dealloc@GOTPCREL]
    movss   xmm0, dword ptr [rsp + 4]

.LBB0_11:
    pop     rax
    ret
```

Now let's change just 1 thing

The type from Vec to array

```
#[inline(never)]  
pub fn sum_arr(arr: [f32; 4]) -> f32 {  
    arr.iter().sum()  
}
```

What actually gets generated now?

That's it

```
example::sum_arr:  
    xorps    xmm0, xmm0  
    addss    xmm0, dword ptr [rdi]  
    addss    xmm0, dword ptr [rdi + 4]  
    addss    xmm0, dword ptr [rdi + 8]  
    addss    xmm0, dword ptr [rdi + 12]  
    ret
```

Static vs Dynamic dispatch

Consider the following simple rust trait:

```
trait Execute {  
    fn execute(&self, x: i32) -> i32;  
}  
  
struct AddX {  
    x: i32,  
}  
  
impl Execute for AddX {  
    fn execute(&self, x: i32) -> i32 {  
        self.x + x  
    }  
}
```

Static vs Dynamic dispatch

Here are the two main ways to execute them:

```
fn execute_dyn(ex: &dyn Execute, x: i32) -> i32 {
    ex.execute(x)
}

fn execute_generic<T: Execute>(ex: &T, x: i32) -> i32 {
    ex.execute(x)
}

pub fn run() {
    let add = AddX { x: 5 };

    dbg!(execute_dyn(&add, 7));
    dbg!(execute_generic(&add, 7));
}
```

Underlying mechanics

LLVM may look intimidating at first

but it's fairly simple

```
define i32 @execute(ptr %self, i32 %x) {  
    // Read the field from the object  
    %_3 = load i32, ptr %self, align 4  
    // Add the field to the argument  
    %_0 = add i32 %_3, %x  
    // Return the result  
    ret i32 %_0  
}  
  
define i32 @execute_dyn(ptr %ex.0, ptr %ex.1, i32 %x) {  
    // Fetch the vtable pointer from the object  
    %0 = getelementptr inbounds ptr, ptr %ex.1, i64 3  
    // Read the function pointer from the vtable  
    %1 = load ptr, ptr %0, align 8  
    // Call the function pointer with the object and the argument  
    %_0 = tail call i32 @%1(ptr %ex.0, i32 %x)  
    // Return the result  
    ret i32 %_0  
}  
  
define internal fastcc i32 @execute_generic(ptr %ex) {  
    // Directly call the execute function with the object and the argument  
    %_0 = tail call i32 @execute(ptr %ex, i32 7)  
    // Return the result  
    ret i32 %_0  
}
```

What's that argument?

It's vtables!

```
@vtable.0 = private unnamed_addr constant <{ ptr, [16 x i8], ptr }>
    <{
        ptr @"_ZN4core3ptr34drop_in_place$LT$example..AddX$GT$17h2205ae1cb6e0e05dE",
        [16 x i8] c"\04\00\00\00\00\00\00\00\04\00\00\00\00\00\00\00",
        ptr @execute
    }>, align 8

define void @run() {
    %add = alloca i32, align 4
    store i32 5, ptr %add, align 4
    %_2 = call i32 @execute_dyn(ptr %add, ptr @vtable.0, i32 7)
    %_21 = call i32 @execute_generic(ptr %add)
    ret void
}
```

So, static or dynamic?

Static:

- Compiler can be smarter about it
- No runtime overhead

Dynamic:

- Easier to do
- Reduces the amount of code generated

Summary

- Tell the compiler your intentions
- The compiler uses those intentions to optimize code
- The compiler can optimize:
 - Data flow
 - Control flow
 - Function calls
 - Algorithms
 - And more

Thanks